

Bewijs Milestone 8

Said khalf

Saif Qudaih

Student 1:

Overzicht vergelijking:

Tabel Rental voor partitionering:

	SEGMENT_NAME	SEGMENT_TYPE	MB	TABLE_COUNT
1	RENTAL	TABLE	56	400000

Stap 2: analyse voor optimalisatie:

```
1 SELECT
2     rentalStatus,
3     COUNT(rentalID) AS total_rentals,
4     ROUND(AVG(totalCost),2) AS avg_cost,
5     MAX(totalCost) AS max_cost,
6     MIN(totalCost) AS min_cost,
7     SUM(totalCost) AS total_revenue
8 FROM Rental
9 WHERE rentalStartDate BETWEEN TO_DATE('2024-01-01', 'YYYY-MM-DD') AND TO_DATE('2024-12-31', 'YYYY-MM-DD')
10    AND totalCost BETWEEN 100 AND 2000
11 GROUP BY rentalStatus
12 ORDER BY total_rentals DESC;
```

Operation	Params	Rows	Total Cost	Raw Desc
⌵ Select		61530	1067.0	cpu_cost = 214046379, io_cost = 10...
⌵ Unknown (PX COORDINATOR)				cpu_cost = null, io_cost = null
⌵ Unknown (PX SEND QC (ORDER))		61530	1067.0	cpu_cost = 214046379, io_cost = 10...
⌵ Order By (SORT ORDER BY)		61530	1067.0	cpu_cost = 214046379, io_cost = 10...
⌵ Unknown (PX RECEIVE)		61530	1067.0	cpu_cost = 214046379, io_cost = 10...
⌵ Unknown (PX SEND RANGE)		61530	1067.0	cpu_cost = 214046379, io_cost = 10...
⌵ [≡] Group By (HASH GROUP BY)		61530	1067.0	cpu_cost = 214046379, io_cost = 10...
⌵ Unknown (PX RECEIVE)		61530	1067.0	cpu_cost = 214046379, io_cost = 10...
⌵ Unknown (PX SEND HASH)		61530	1067.0	cpu_cost = 214046379, io_cost = 10...
⌵ [≡] Group By (HASH GROUP BY)		61530	1067.0	cpu_cost = 214046379, io_cost = 10...
⌵ Unknown (PX BLOCK ITERATOR)		61530	1064.0	cpu_cost = 96409022, io_cost = 1061
Full Scan (TABLE ACCESS FULL table: RENTAL;		61530	1064.0	cpu_cost = 96409022, io_cost = 1061

Bewijs Milestone 8

Said khalf

Saif Qudaih

NA partitioning:

```
5
6 CREATE TABLE Rental (
7     rentalID NUMBER GENERATED ALWAYS AS IDENTITY,
8     chassisNr VARCHAR2(50),
9     customerID INTEGER NOT NULL,
10    rentalStartDate DATE NOT NULL,
11    rentalEndDate DATE NOT NULL,
12    rentalStatus VARCHAR2(20) NOT NULL,
13    paymentDetails VARCHAR2(200),
14    customerFeedback VARCHAR2(500),
15    insuranceDetails VARCHAR2(200),
16    totalCost DECIMAL(10, 2) NOT NULL,
17    shopID INTEGER NOT NULL,
18    CONSTRAINT RENTAL_PK PRIMARY KEY (rentalID),
19    CONSTRAINT RENTAL_FK_chassisNr FOREIGN KEY (chassisNr)
20        REFERENCES MotorBike (chassisNr),
21    CONSTRAINT RENTAL_FK_CUSTOMERID FOREIGN KEY (customerID)
22        REFERENCES Customer (customerID),
23    CONSTRAINT RENTAL_FK_SHOPID FOREIGN KEY (shopID)
24        REFERENCES Shop (shopID),
25    CONSTRAINT UNIQUE_RENTAL_COMBINATION UNIQUE (rentalID, customerID, chassisNr, rentalStartDate)
26 )
27 PARTITION BY RANGE (rentalStartDate)
28 INTERVAL (NUMTOYMINTERVAL(1, 'MONTH'))
29 (
30     PARTITION p202301 VALUES LESS THAN (TO_DATE('2023-02-01', 'YYYY-MM-DD')),
31     PARTITION p202302 VALUES LESS THAN (TO_DATE('2023-03-01', 'YYYY-MM-DD')),
32     PARTITION p202303 VALUES LESS THAN (TO_DATE('2023-04-01', 'YYYY-MM-DD'))
33 );
34
```

Tabel Rental NA partitioning:

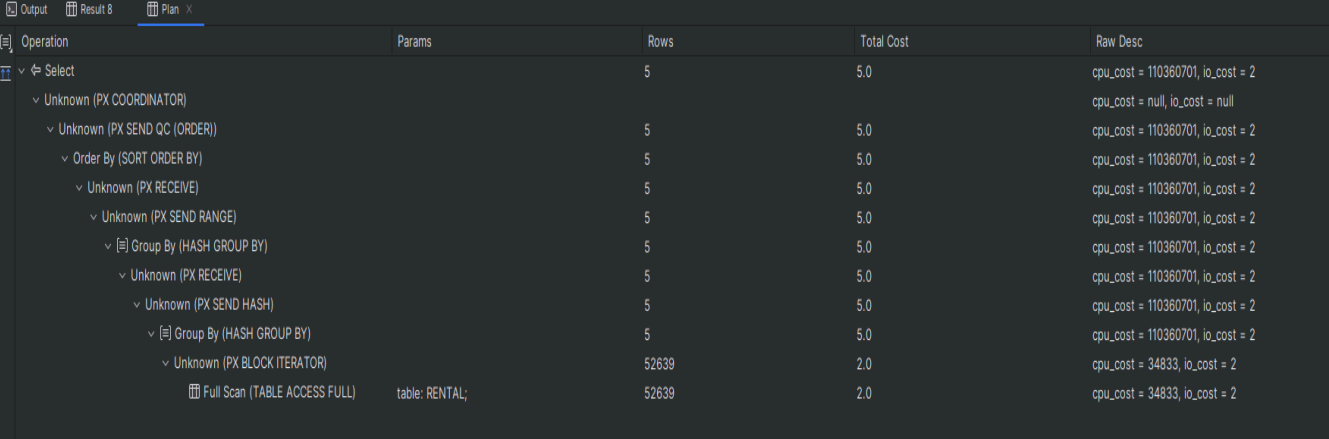
Output Result 8 x Plan			
1 row v			
SEGMENT_NAME	SEGMENT_TYPE	MB	TABLE_COUNT
1 RENTAL	TABLE PARTITION	104	400000

Bewijs Milestone 8

Said khalf

Saif Qudaih

Explain plan na partitionering



Operation	Params	Rows	Total Cost	Raw Desc
Select		5	5.0	cpu_cost = 110360701, io_cost = 2
Unknown (PX COORDINATOR)				cpu_cost = null, io_cost = null
Unknown (PX SEND QC (ORDER))		5	5.0	cpu_cost = 110360701, io_cost = 2
Order By (SORT ORDER BY)		5	5.0	cpu_cost = 110360701, io_cost = 2
Unknown (PX RECEIVE)		5	5.0	cpu_cost = 110360701, io_cost = 2
Unknown (PX SEND RANGE)		5	5.0	cpu_cost = 110360701, io_cost = 2
Group By (HASH GROUP BY)		5	5.0	cpu_cost = 110360701, io_cost = 2
Unknown (PX RECEIVE)		5	5.0	cpu_cost = 110360701, io_cost = 2
Unknown (PX SEND HASH)		5	5.0	cpu_cost = 110360701, io_cost = 2
Group By (HASH GROUP BY)		5	5.0	cpu_cost = 110360701, io_cost = 2
Unknown (PX BLOCK ITERATOR)		52639	2.0	cpu_cost = 34833, io_cost = 2
Full Scan (TABLE ACCESS FULL)	table: RENTAL;	52639	2.0	cpu_cost = 34833, io_cost = 2

Conclusie:

De partitionering van de Rental-tabel heeft de query prestaties aanzienlijk verbeterd door de totale kosten te verlagen van 1.067,0 naar 5,0. Het aantal geraadpleegde rijen is ook afgenomen van 61.530 naar 52.639, wat duidt op efficiëntere gegevensopvraging.

Dit geeft aan dat partitionering op rentalStartDate de query heeft geoptimaliseerd, waarschijnlijk door een betere afstemming met de filter voorwaarden van de query en efficiëntere gegevens scanning.

De substantiële vermindering van de totale kosten suggereert dat partition pruning effectief het gescande gegevensvolume beperkt, wat bijdraagt aan verbeterde query prestaties.

Bewijs Milestone 8

Said khalf

Saif Qudaih

Student 2

Voor Materialized View:

	SEGMENT_NAME	SEGMENT_TYPE	MB	TABLE_COUNT
1	SERVICE	TABLE	43	540000

Queries:

```
-- ANALYTISCHE QUERY
-- Deze query geeft de voornaam en achternaam van elke medewerker terug,
-- samen met het aantal services dat ze hebben.
SELECT e.firstName, e.lastName, COUNT(s.serviceld) AS Aantal_Services
FROM SERVICE s
      JOIN Employee e ON s.EMPLOYEEID = e.EMPLOYEEID
      JOIN Department d ON e.departmentID != d.departmentID
GROUP BY e.firstName, e.lastName
ORDER BY Aantal_Services DESC;
```

Explain plan:

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	TQ	IN-OUT	PQ Distrib
0	SELECT STATEMENT		26100	4078K	862 (2)	00:00:01			
1	PX COORDINATOR								
2	PX SEND QC (ORDER)	:TQ10004	26100	4078K	862 (2)	00:00:01	Q1,04	P->S	QC (ORDER)
3	SORT ORDER BY		26100	4078K	862 (2)	00:00:01	Q1,04	PCWP	
4	PX RECEIVE		26100	4078K	862 (2)	00:00:01	Q1,04	PCWP	
5	PX SEND RANGE	:TQ10003	26100	4078K	862 (2)	00:00:01	Q1,03	P->P	RANGE
6	HASH GROUP BY		26100	4078K	862 (2)	00:00:01	Q1,03	PCWP	
7	PX RECEIVE		26100	4078K	862 (2)	00:00:01	Q1,03	PCWP	
8	PX SEND HASH	:TQ10002	26100	4078K	862 (2)	00:00:01	Q1,02	P->P	HASH
9	HASH GROUP BY		26100	4078K	862 (2)	00:00:01	Q1,02	PCWP	
*10	HASH JOIN		26100	4078K	860 (2)	00:00:01	Q1,02	PCWP	
11	PX RECEIVE		900	114K	3 (0)	00:00:01	Q1,02	PCWP	
12	PX SEND BROADCAST	:TQ10000	900	114K	3 (0)	00:00:01	Q1,00	P->P	BROADCAST
13	PX BLOCK ITERATOR		900	114K	3 (0)	00:00:01	Q1,00	PCWC	
14	TABLE ACCESS FULL	EMPLOYEE	900	114K	3 (0)	00:00:01	Q1,00	PCWP	
15	MERGE JOIN CARTESIAN		27540	806K	857 (2)	00:00:01	Q1,02	PCWP	
16	TABLE ACCESS FULL	DEPARTMENT	30	390	2 (0)	00:00:01	Q1,02	PCWP	
17	BUFFER SORT		918	15606	855 (2)	00:00:01	Q1,02	PCWP	
18	VIEW	VW_GBF_7	918	15606			Q1,02	PCWP	
19	HASH GROUP BY		918	3672	825 (2)	00:00:01	Q1,02	PCWP	
20	PX RECEIVE		918	3672	825 (2)	00:00:01	Q1,02	PCWP	
21	PX SEND HASH	:TQ10001	918	3672	825 (2)	00:00:01	Q1,01	P->P	HASH
22	HASH GROUP BY		918	3672	825 (2)	00:00:01	Q1,01	PCWP	
23	PX BLOCK ITERATOR		540K	2109K	817 (1)	00:00:01	Q1,01	PCWC	
24	TABLE ACCESS FULL	SERVICE	540K	2109K	817 (1)	00:00:01	Q1,01	PCWP	

Bewijs Milestone 8

Said khalf

Saif Qudaih

Explain Plan na Materialized View:

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	TQ	IN-OUT	PQ Distrib
0	SELECT STATEMENT		100	1900	2 (0)	00:00:01			
1	PX COORDINATOR								
2	PX SEND QC (RANDOM)	:TQ10000	100	1900	2 (0)	00:00:01	Q1,00	P->S	QC (RAND)
3	PX BLOCK ITERATOR		100	1900	2 (0)	00:00:01	Q1,00	PCWC	
4	MAT_VIEW ACCESS FULL	MOST_SERVICES_PER_EMPLOYEE_MV	100	1900	2 (0)	00:00:01	Q1,00	PCWP	

Na nieuwe data:

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	TQ	IN-OUT	PQ Distrib
0	SELECT STATEMENT		100	1900	2 (0)	00:00:01			
1	PX COORDINATOR								
2	PX SEND QC (RANDOM)	:TQ10000	100	1900	2 (0)	00:00:01	Q1,00	P->S	QC (RAND)
3	PX BLOCK ITERATOR		100	1900	2 (0)	00:00:01	Q1,00	PCWC	
4	MAT_VIEW ACCESS FULL	MOST_SERVICES_PER_EMPLOYEE_MV	100	1900	2 (0)	00:00:01	Q1,00	PCWP	

Conclusie:

Als je kunt vergelijken de eerste Explain plan met de tweede, dan kunt je zien dat de cost is veel minder. Dus in de materialized view die ik heb gebruikt bewijst dat de cost is minder.

En toen ik mijn database opnieuw vulde of extra data toevoegde, bleef het resultaat van de materialized view hetzelfde als ervoor.

Een nadeel van een materialized view is dat deze gerefreshed moet worden wanneer de originele data verandert. Als de data vaak verandert, is een materialized view geen goede keuze.